


Foundations of Statistical Inference

...

18 February 2019
PLSC 309

Review: signal or noise?




Review: signal or noise?

After opening a new school, a city wants to determine if the new school is performing well. The average school in the district has a graduation rate of 74%, with a standard deviation of 8%. The new school's graduation rate is 82%.

Review: signal or noise?

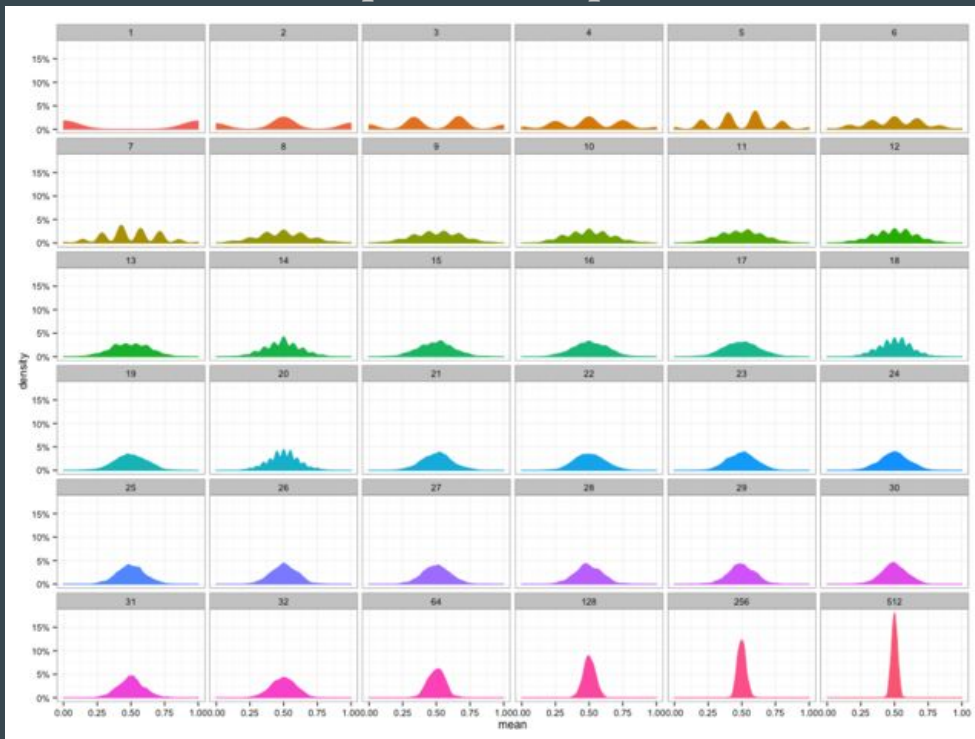
After opening a new school, a city wants to determine if the new school is performing well. The average school in the district has a graduation rate of 74%, with a standard deviation of 8%. The new school's graduation rate is 82%. **Is this increase due to random chance, or has the new school done something different than the rest of the other schools to improve performance?**

Review: sampling distributions



Review: CLT and normal approximation


For all i.i.d. Variables, the means of repeated samples of sufficient size will be normally distributed.




All distributions have parameters

- Binomial: p (probability)
- Poisson: λ (average count)
- Student's T: ν (normality parameter)
- Normal: μ, σ (mean, standard deviation)


Parameters control the shape and location of a distribution




Parameters control the shape and location of a distribution



Parameters control the shape and location of a distribution



Parameters control the shape and location of a distribution



Parameters are probabilistic

- They are characteristics of infinite processes
- All observations in the universe
- We estimate the value of a parameter with a *point estimate*
- Population = parameter
- Sample = point estimate

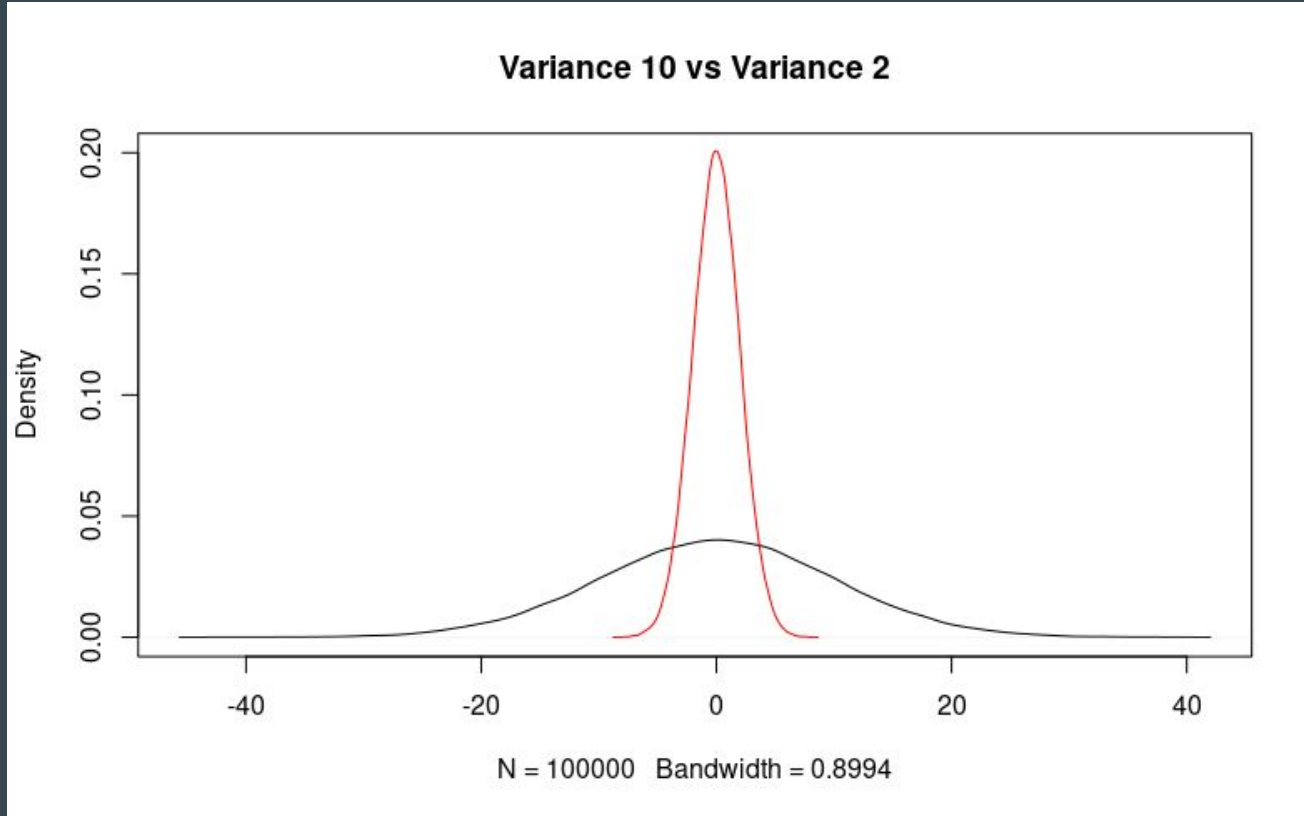
Some point estimates

- Binomial: p (probability)
 - $P = k / n$
- Poisson: λ (average count)
 - $\lambda = \text{mean}(X)$
- Normal: μ, σ (mean, standard deviation)
 - μ : `mean(x)`
 - σ : `sqrt(var(x))`

Point estimate examples


- Survey of 1000 likely voters finds 41% approval for President Trump
 - Point estimate: mean(approval) for sample
 - Parameter: average support amongst entire population
- The variance of test scores in a particular school is 8.72
 - Point estimate: variance for that school
 - Parameter: overall variance of all schools in the district

Uncertainty in point estimates




Uncertainty in point estimates

10 data points randomly drawn from standard normal



10,000 data points randomly drawn from standard normal



Uncertainty in point estimates

Uncertainty comes from two places:

- Variance
- Sample size

Standard Error

- Standard error captures the *uncertainty* of point estimates

$$SE = \frac{\sigma}{\sqrt{n}}$$

Standard Error

A sample of 1000 likely voters finds 62% support legalizing marijuana, with a variance of 4%.

$$SE = \sqrt{4\%} / \sqrt{1000} = 6.3\%$$


What do we do with standard errors?

- Standard errors provide a measure of uncertainty for our point estimates
- We can compare the quality of point estimates on the same scale...
- ...but not really helpful if we're only looking at a single point


Confidence intervals

- Confidence intervals also measure uncertainty, like standard errors
- They translate standard errors from being in units of X to units of Y (i.e. probability)
- We want to know the *probability* that our point estimate is within a certain range of values

Point estimates follow a normal distribution



Point estimates follow a normal distribution



Confidence interval formula

- Select z-score for probability level you're interested in
 - $Z = 1$; 68% confident
 - $Z = 2$; 95% confident
 - $Z = 3$; 99.7% confident
- $CI = \text{point estimate} \mp z * SE$

Confidence Interval example

A sample of 1000 likely voters finds 62% support legalizing marijuana, with a variance of 4%.

$$SE = \sqrt{4\%} / \sqrt{1000} = 6.3\%$$

$$95\% \text{ CI} = 62\% \pm 2 * 6.3\%$$

$$95\% \text{ CI} = 49.4\% - 74.6\%$$

Confidence interval interpretation

A confidence interval means that there is 95% chance the true population parameter lies within the range of the C.I.

- In other words, 95 out of 100 perfectly random samples will have C.I. that overlap with the true population mean
- All based on CLT, so all CLT assumptions apply

Null Hypothesis Significance Testing: overview

The null hypothesis states that nothing changes.

H_0 : Nothing changes

- Any deviation from $E(X)$ due to random chance

H_A : Something changes

- Deviation from $E(X)$ due to systematic change


Null Hypothesis Significance Testing: overview

The null hypothesis states that $p(x_i = \text{point estimate}) \square \text{null distribution}$

$$H_0: p(x_i = \text{point estimate}) \square \text{null distribution}$$

$$H_A: p(x_i = \text{point estimate}) \square \text{alternative distribution}$$


Null distribution: nothing changes



α : critical value

- This is the level at which you deem an observation “extreme” or “out of the ordinary”
- In other words, it’s so extreme, that we actually think it is drawn from a different distribution than the rest of our data
- α is often set to the fifth percentile of extreme values

α : critical value




NHST steps

1. Calculate z-score for point estimate you're interested in
 - Use standard error instead of standard deviation
 - $P(x_i) = P(z_i)$
2. Evaluate z-score with respect to the null distribution
 - $\text{Normal}(\mu=0, \sigma=1)$
3. If $p(z_i < \alpha)$ is true
 - Reject H_0 / Accept H_A
4. If $p(z_i > \alpha)$ is true
 - Accept H_0 / Reject H_A

$p(z_i) = \text{P-value}$

- If the null distribution is true, we have an α chance to get the point estimate that we got
- If the p-value is below α we say that we accept the alternative hypothesis
 - Our point estimate is so rare in the null distribution, that we think it follows some alternative distribution

Null vs alternative distributions



NHST example

After a presidential debate, the candidate has an approval rating, drawn from a sample of 100 voters, of 53%. The average of past approval polls has been 48% with a 0.5% standard deviation. Did the debate make a difference in the candidate's approval?

1. Calculate z-score

After a presidential debate, the candidate has an approval rating, drawn from a sample of 100 voters, of 53%. The average of past approval polls has been 48% with a 0.5% standard deviation. Did the debate make a difference in the candidate's approval?

$$Z = \frac{\bar{X} - \mu}{\frac{\sigma}{\sqrt{N}}}$$

where,

\bar{X} = the sample mean

μ = the population mean

σ = the population standard deviation

N = the sample size

1. Calculate z-score

After a presidential debate, the candidate has an approval rating, drawn from a sample of 100 voters, of 53%. The average of past approval polls has been 48% with a 5% standard deviation. Did the debate make a difference in the candidate's approval?

$$Z = 53 - 48 / (5/\text{sqrt}(100)) = 2.5$$

2. Evaluate with respect to the null distribution

After a presidential debate, the candidate has an approval rating, drawn from a sample of 100 voters, of 53%. The average of past approval polls has been 48% with a 5% standard deviation. Did the debate make a difference in the candidate's approval?

$$Z_i = 2.5$$

$$P(z_i = 2.5 \mid \text{null distribution}) = .006$$

3. If $p(z_i < \alpha)$, reject H_0 and accept H_A

After a presidential debate, the candidate has an approval rating, drawn from a sample of 100 voters, of 53%. The average of past approval polls has been 48% with a 5% standard deviation. Did the debate make a difference in the candidate's approval?

We would expect to see a 53% approval rating, assuming nothing changed, only 0.6% of the time. Therefore, we argue that the debate changed the underlying distribution of voter approval, and our 53% approval rating was drawn from this alternative distribution.

We can also do NHST through confidence intervals

After a presidential debate, the candidate has an approval rating, drawn from a sample of 100 voters, of 53%. The average of past approval polls has been 48% with a 5% standard deviation. Did the debate make a difference in the candidate's approval?

Point estimate $\pm z * SE$

$$53 \pm 2 * (5 / \sqrt{100}) = (52\%, 54\%)$$

(52, 54) does not overlap with 48%

What p-values do not do

1. P-values are not statements about whether the alternative hypothesis is true, or whether the null hypothesis is false
2. P-values are not the probability that a given event is the product of random chance
3. Smaller p-values do not necessarily mean larger effects

P-values are not a “truth percentage”

- A p-value is NOT the % chance that the alternative distribution is true!
- The only reason we know the null distribution is because of CLT
 - Main assumption of CLT is i.i.d.
- We have no idea what the alternative distribution is
- P-values can only be interpreted with respect to the null distribution

P-values are not likelihood something due to random chance

- Say we calculate a .04 p-value
- That does not mean there is a 4% chance this is due to random chance
- NHST *assumes null distribution is TRUE*
 - *Assumes that all variation is purely due to random chance*
- Calculates the likelihood of a certain point estimate under that assumption
- This tells us nothing about whether that assumption is true!

P-values are not about size of effect

After a presidential debate, the candidate has an approval rating, drawn from a sample of 100 voters, of 53%. The average of past approval polls has been 48% with a 5% standard deviation. Did the debate make a difference in the candidate's approval?

P-values are not about size of effect

After a presidential debate, the candidate has an approval rating, drawn from a sample of **1000 voters**, of 53%. The average of past approval polls has been 48% with a 5% standard deviation. Did the debate make a difference in the candidate's approval?

$$Z = 53 - 48 / (5/\text{sqrt}(1000)) = 31$$

$$\text{P-value} = .0000000000003$$

Review

- We want to estimate parameters of distributions, like the mean
 - We call these “point estimates”
- We use standard errors to quantify our uncertainty about those point estimates
- Using the CLT, we can construct confidence intervals to provide a range of plausible values
- We can also use the CLT to do Null Hypothesis Significance Testing
 - Assume the null distribution
 - Provide a probability (p-value) of the likelihood of a given point estimate under the null
 - Reject or accept a value as too extreme to suggest a different distribution